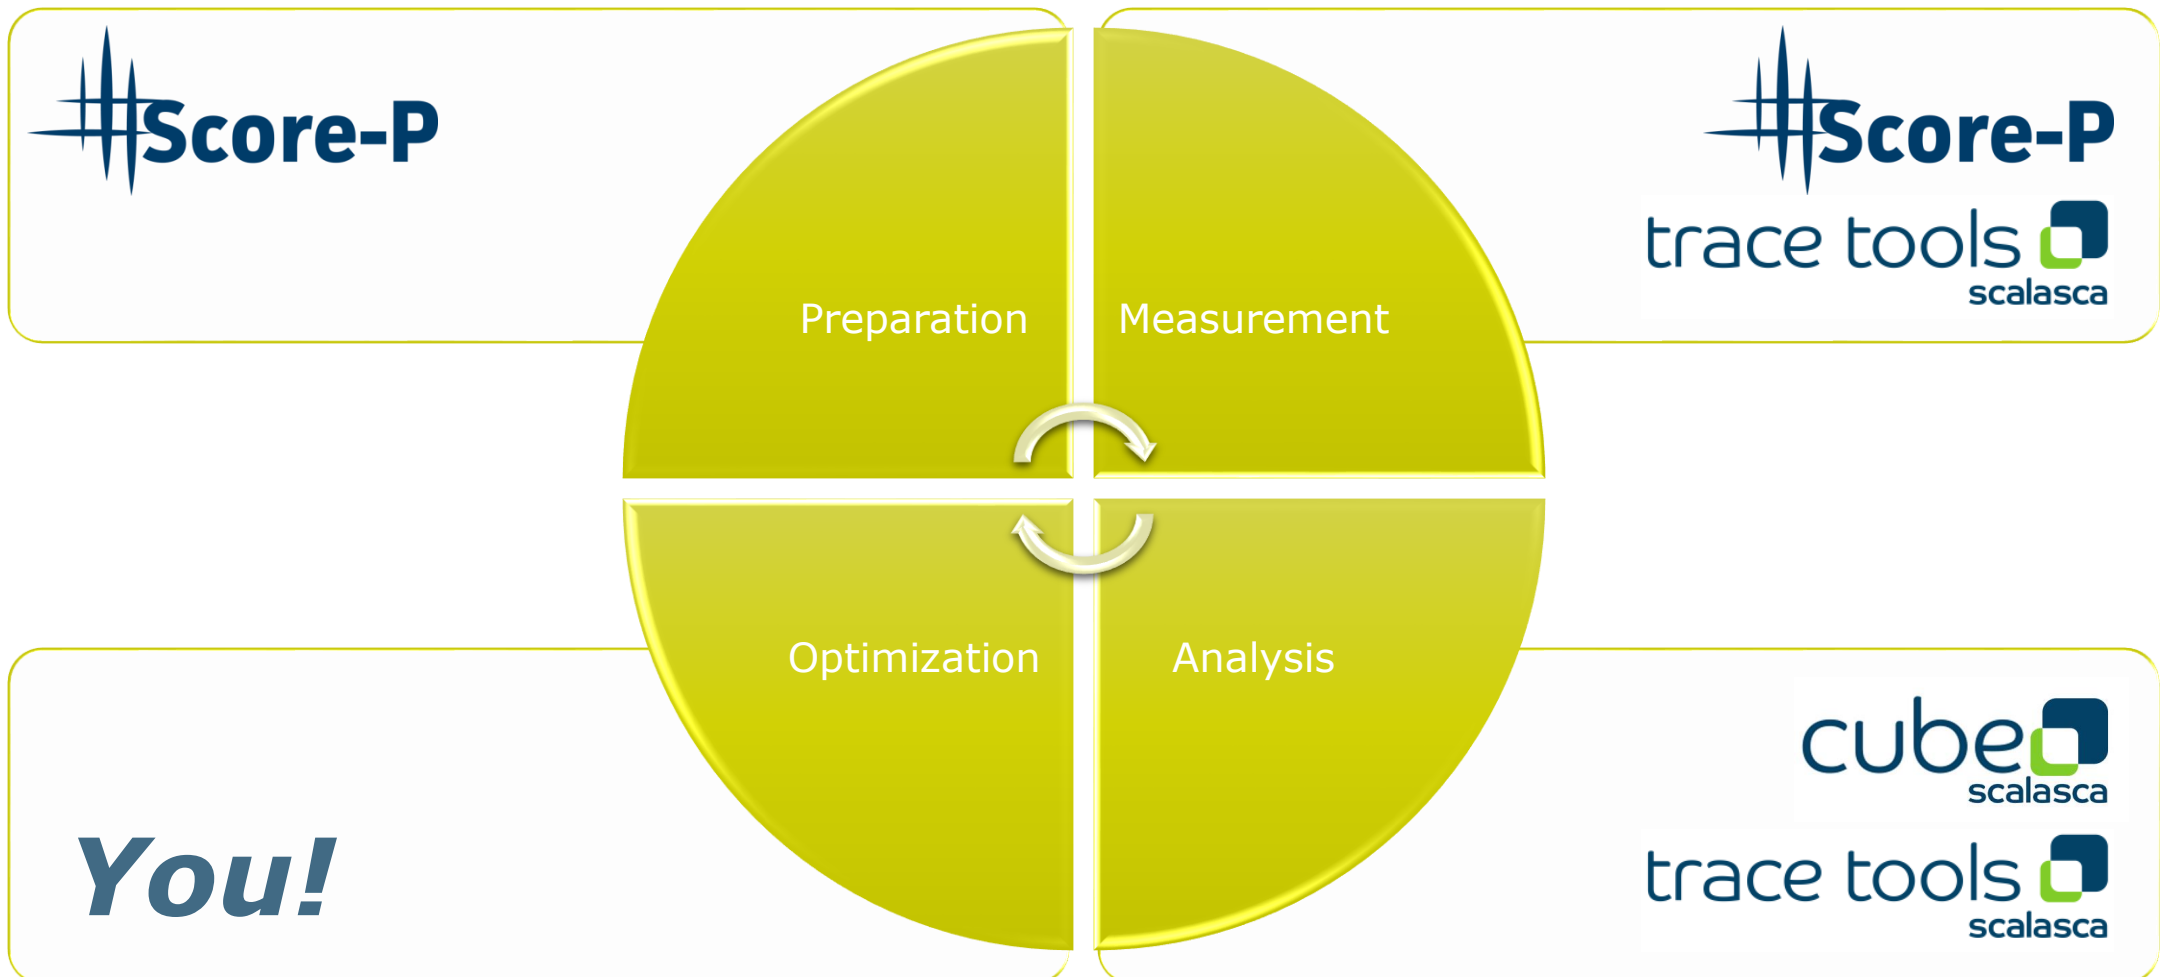


Performance Analysis Workshop Series

Session 4: Score-P / Cube / Scalasca

Markus Geimer
Brian Wylie
Jülich Supercomputing Centre

Performance engineering workflow





- Infrastructure for instrumentation and performance measurements
- Instrumented application can be used to produce several results:
 - Call-path profiling: CUBE4 data format used for data exchange
 - Event-based tracing: OTF2 data format used for data exchange
- Supported parallel paradigms:
 - Multi-process: MPI, SHMEM
 - Thread-parallel: OpenMP, POSIX threads
 - Accelerator-based: CUDA, HIP, OpenCL, OpenACC

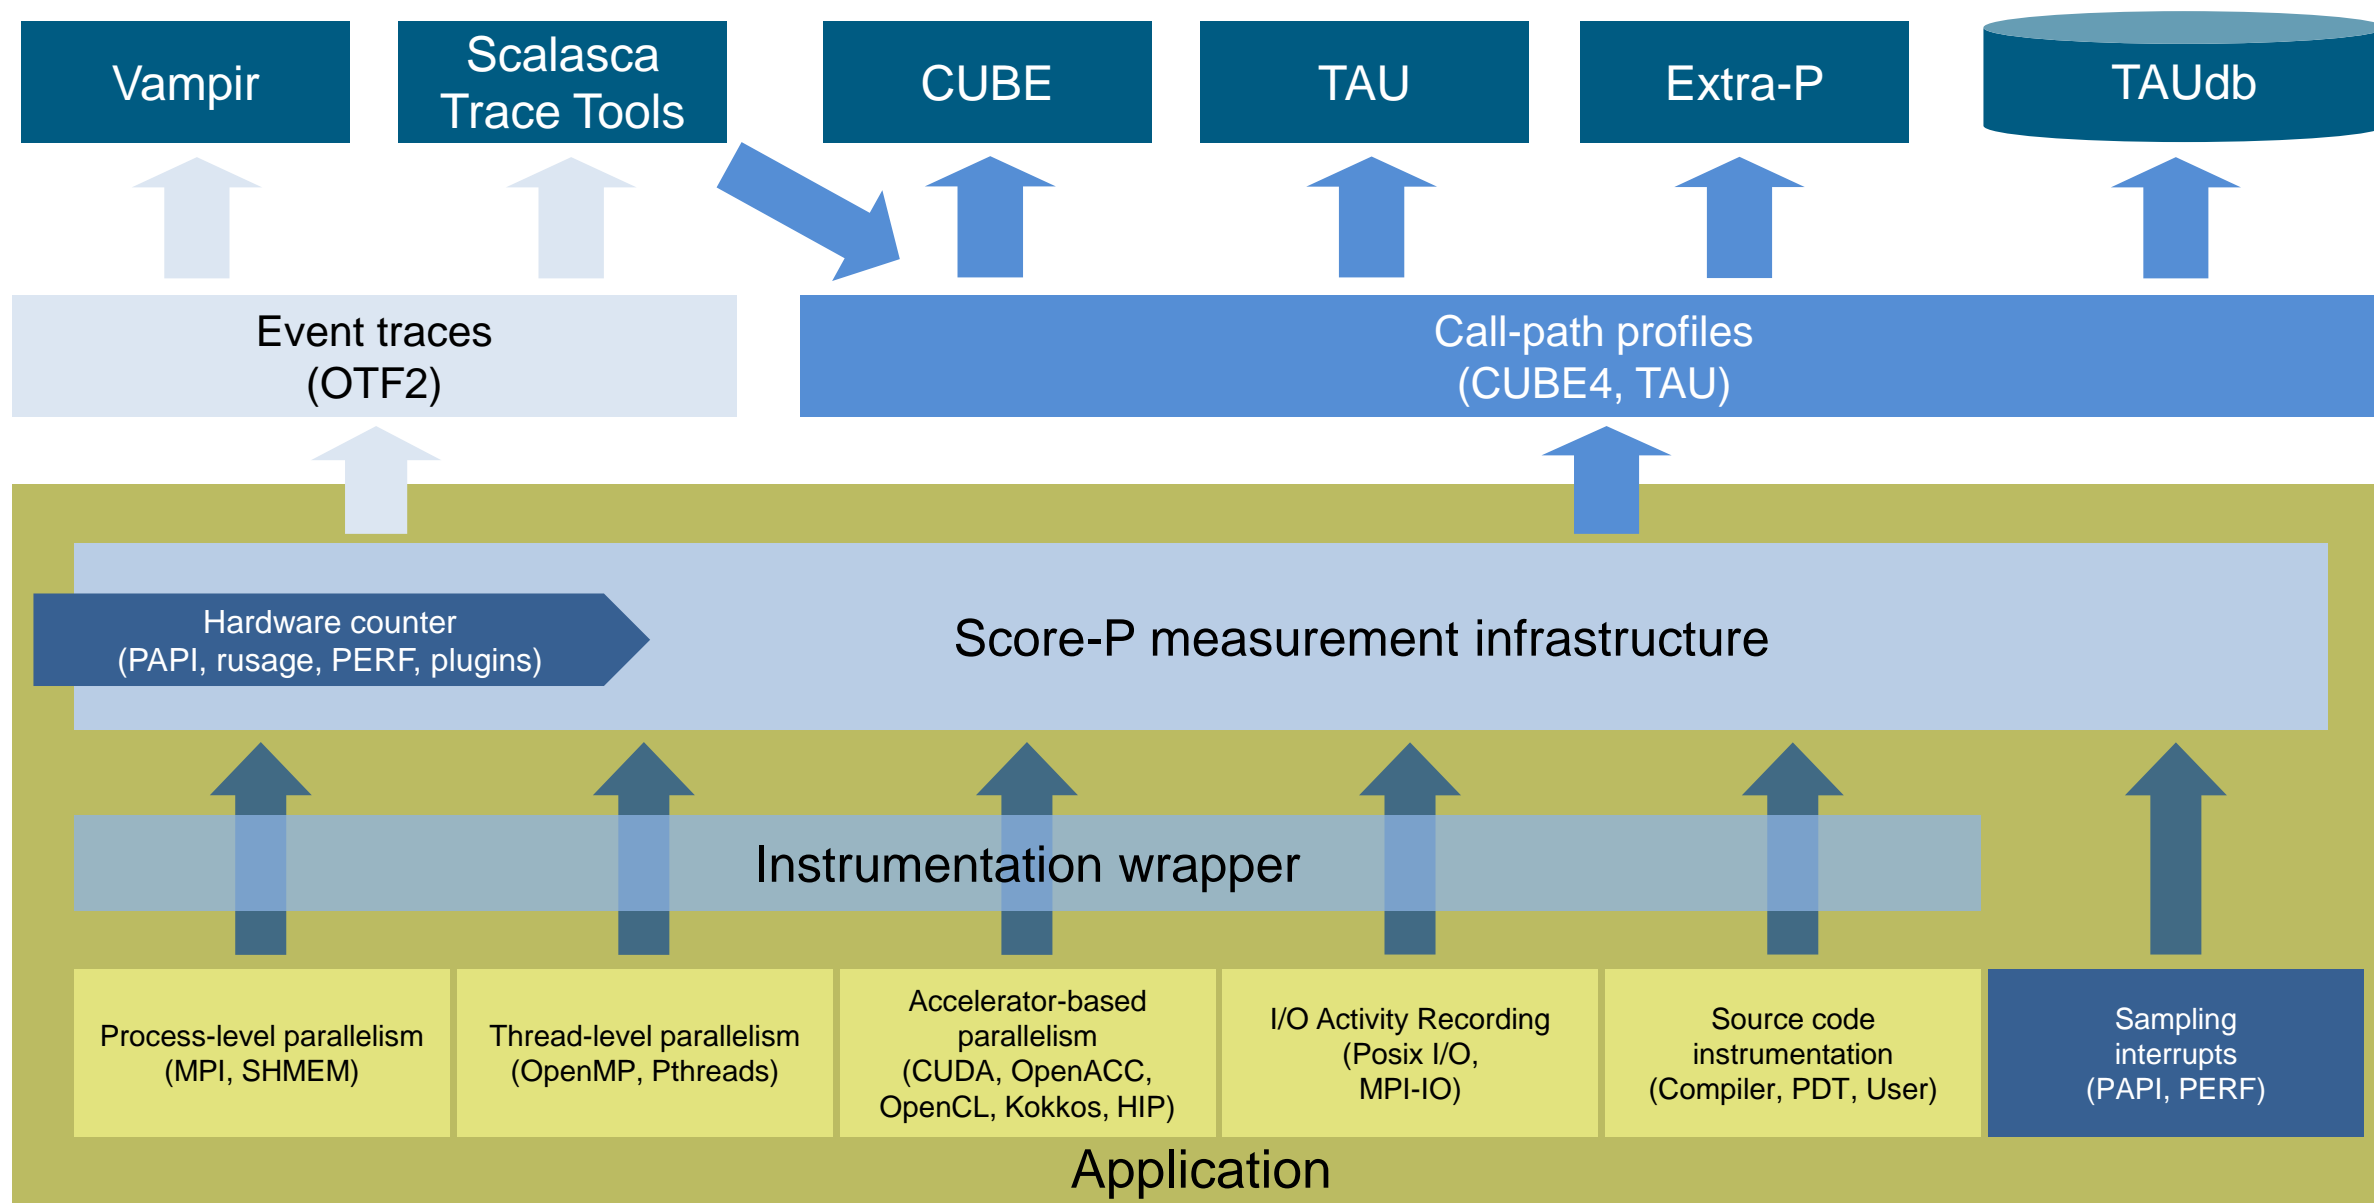
- Initial project funded by BMBF
- Close collaboration with PRIMA project funded by DOE
- Further developed in multiple 3rd-party funded projects

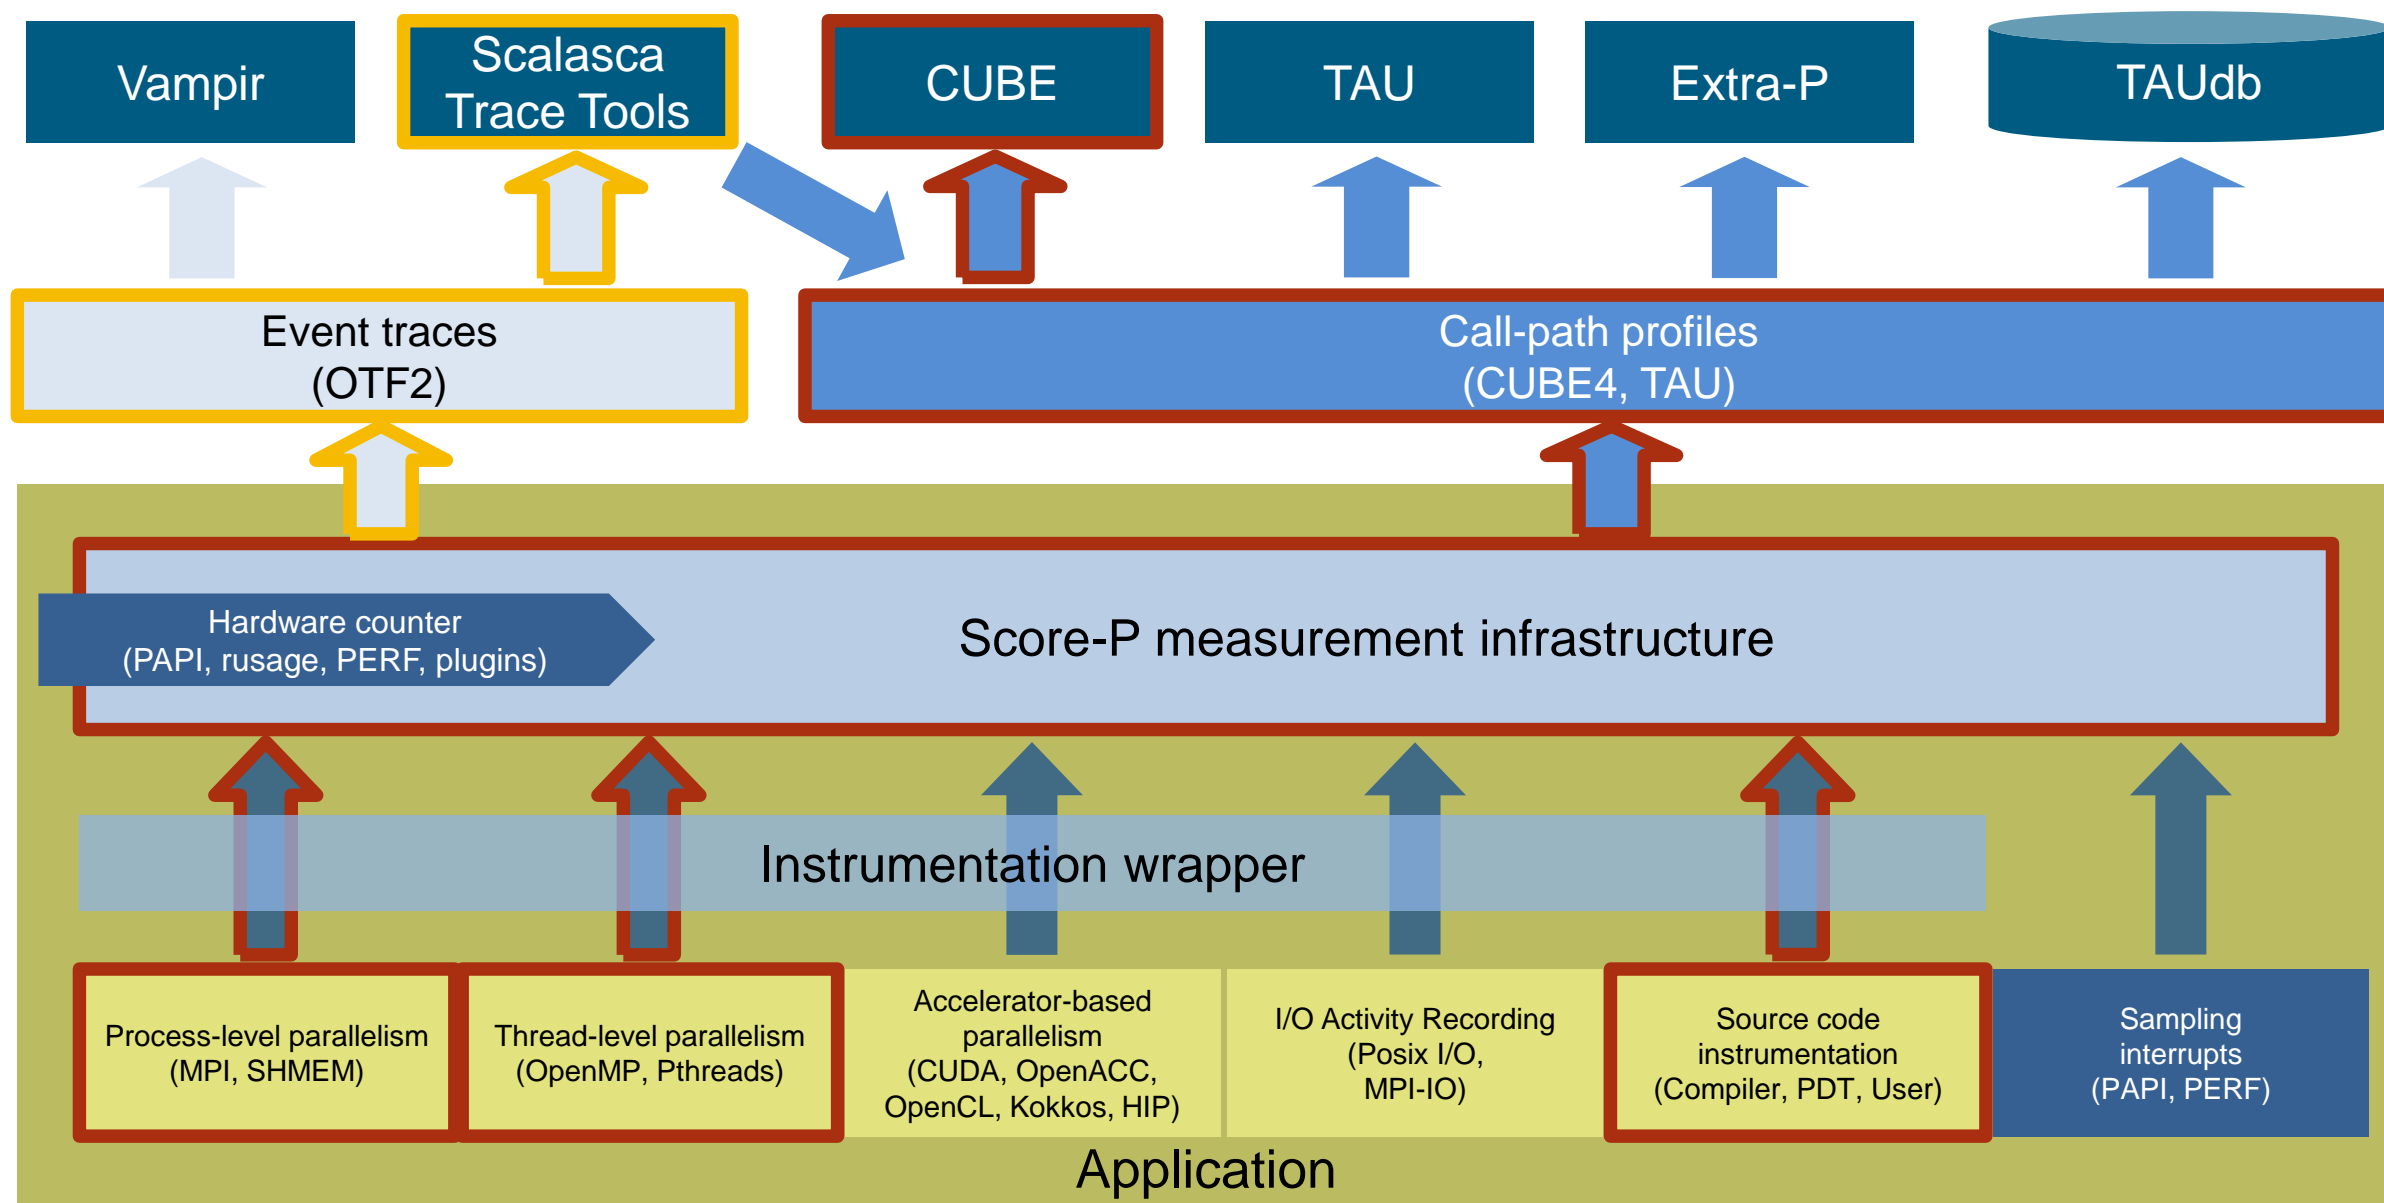
GEFÖRDERT VOM



Bundesministerium
für Bildung
und Forschung







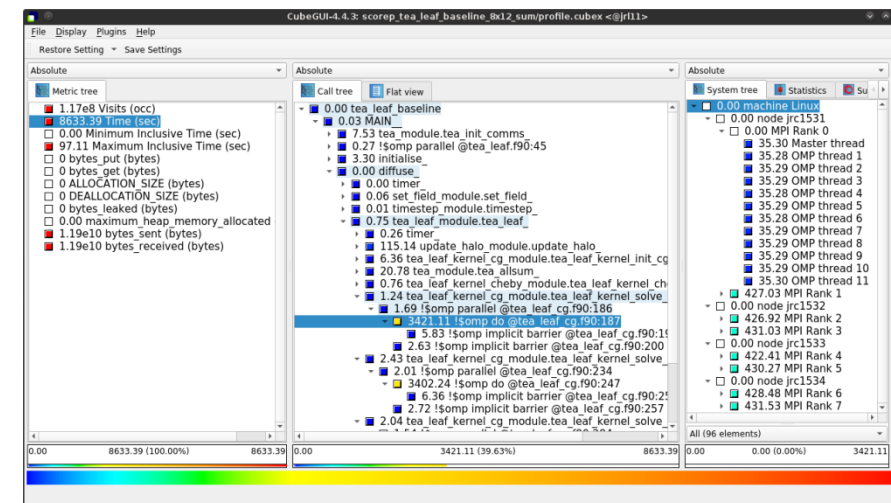
Score-P features

- Open source: 3-clause BSD license
 - Commitment to joint long-term cooperation
 - Development based on meritocratic governance model
 - Open for contributions and new partners
- Portability: supports all major HPC platforms
- Scalability: successful measurements with >1M threads
- Functionality:
 - Generation of call-path profiles and event traces (supporting highly scalable I/O)
 - Using direct instrumentation and sampling
 - Flexible measurement configuration without re-compilation
 - Recording of time, visits, communication data, hardware counters
 - Support for MPI, SHMEM, OpenMP, Pthreads, CUDA, HIP, OpenCL, OpenACC and valid combinations
- Latest release: Score-P 8.4 (Mar 2024)



CubeLib DOI 10.5281/zenodo.1248078
CubeGUI DOI 10.5281/zenodo.1248087


- Parallel program analysis report exploration tools
 - Libraries for XML+binary report reading & writing
 - Algebra utilities for report processing
 - GUI for interactive analysis exploration
 - Requires Qt ≥ 5
- Originally developed as part of the Scalasca toolset
- Now available as separate components
 - Can be installed independently of Score-P and Scalasca, e.g., on laptop/desktop
 - Latest releases: Cube v4.8.2 (Sep 2023)



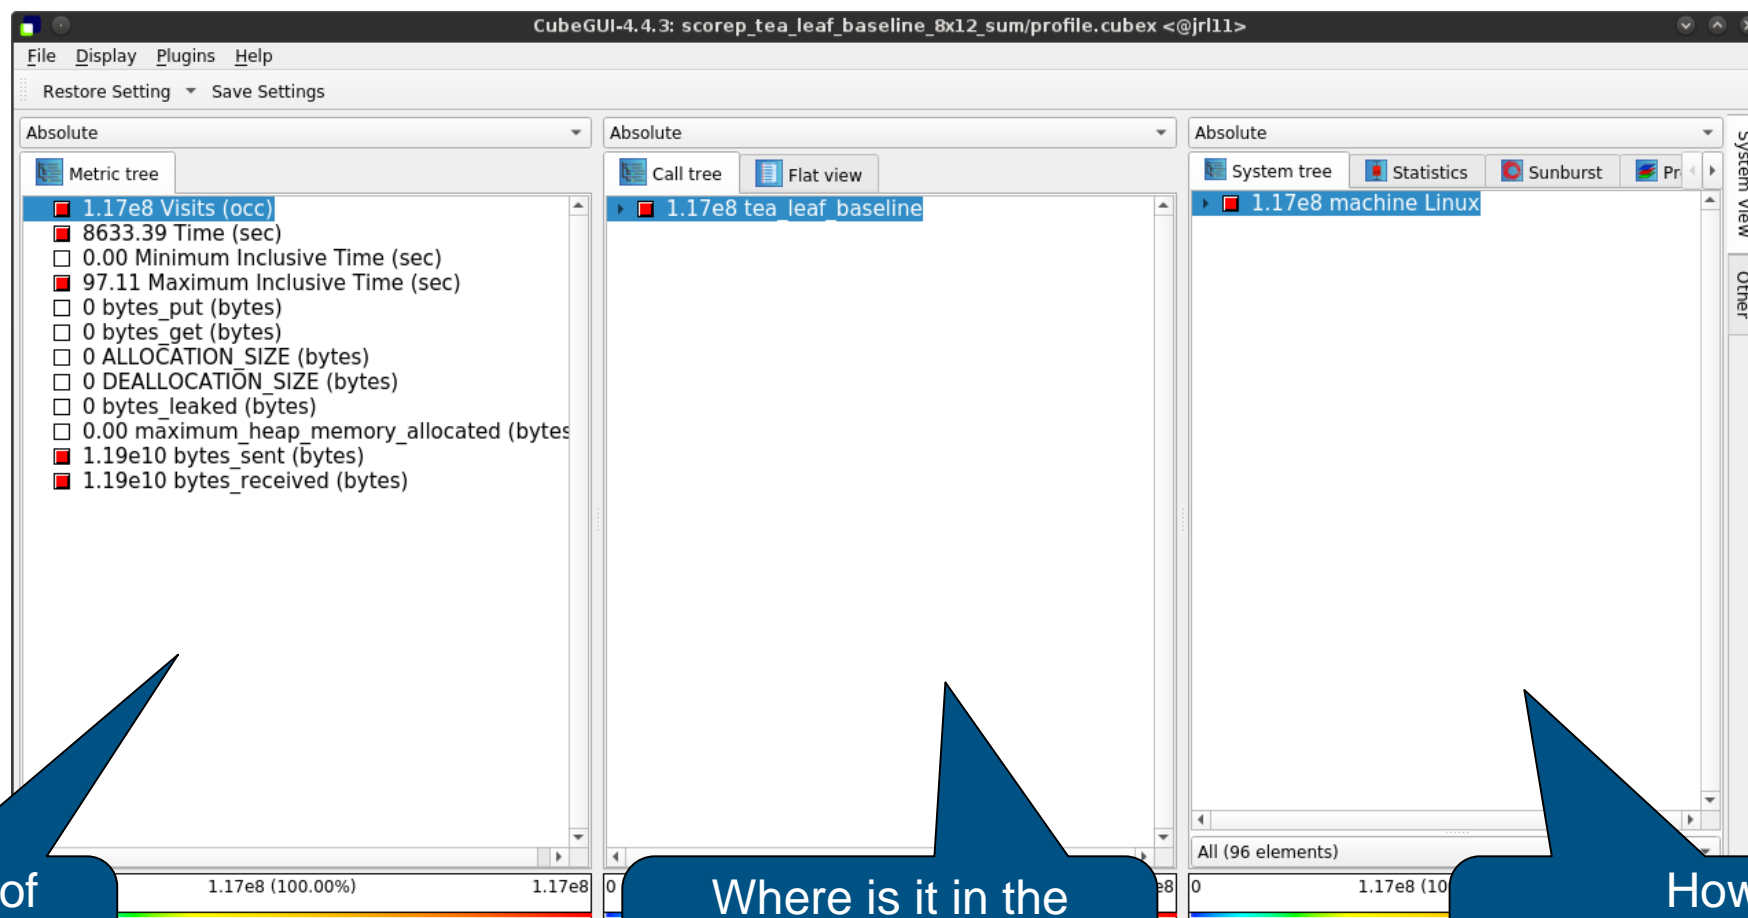
Note: source distribution tarballs for Linux, as well as binary packages provided for Linux, Windows & MacOS, from www.scalasca.org website in Software/Cube 4.x

Analysis presentation and exploration

- Representation of values (severity matrix) on three hierarchical axes
 - Performance property (metric)
 - Call path (program location)
 - System location (process/thread)
- Three coupled tree browsers
- Cube displays severities
 - As value: for precise comparison
 - As color: for easy identification of hotspots
 - Inclusive value when closed & exclusive value when expanded
 - Customizable via display modes



Plain summary analysis report (opening view)

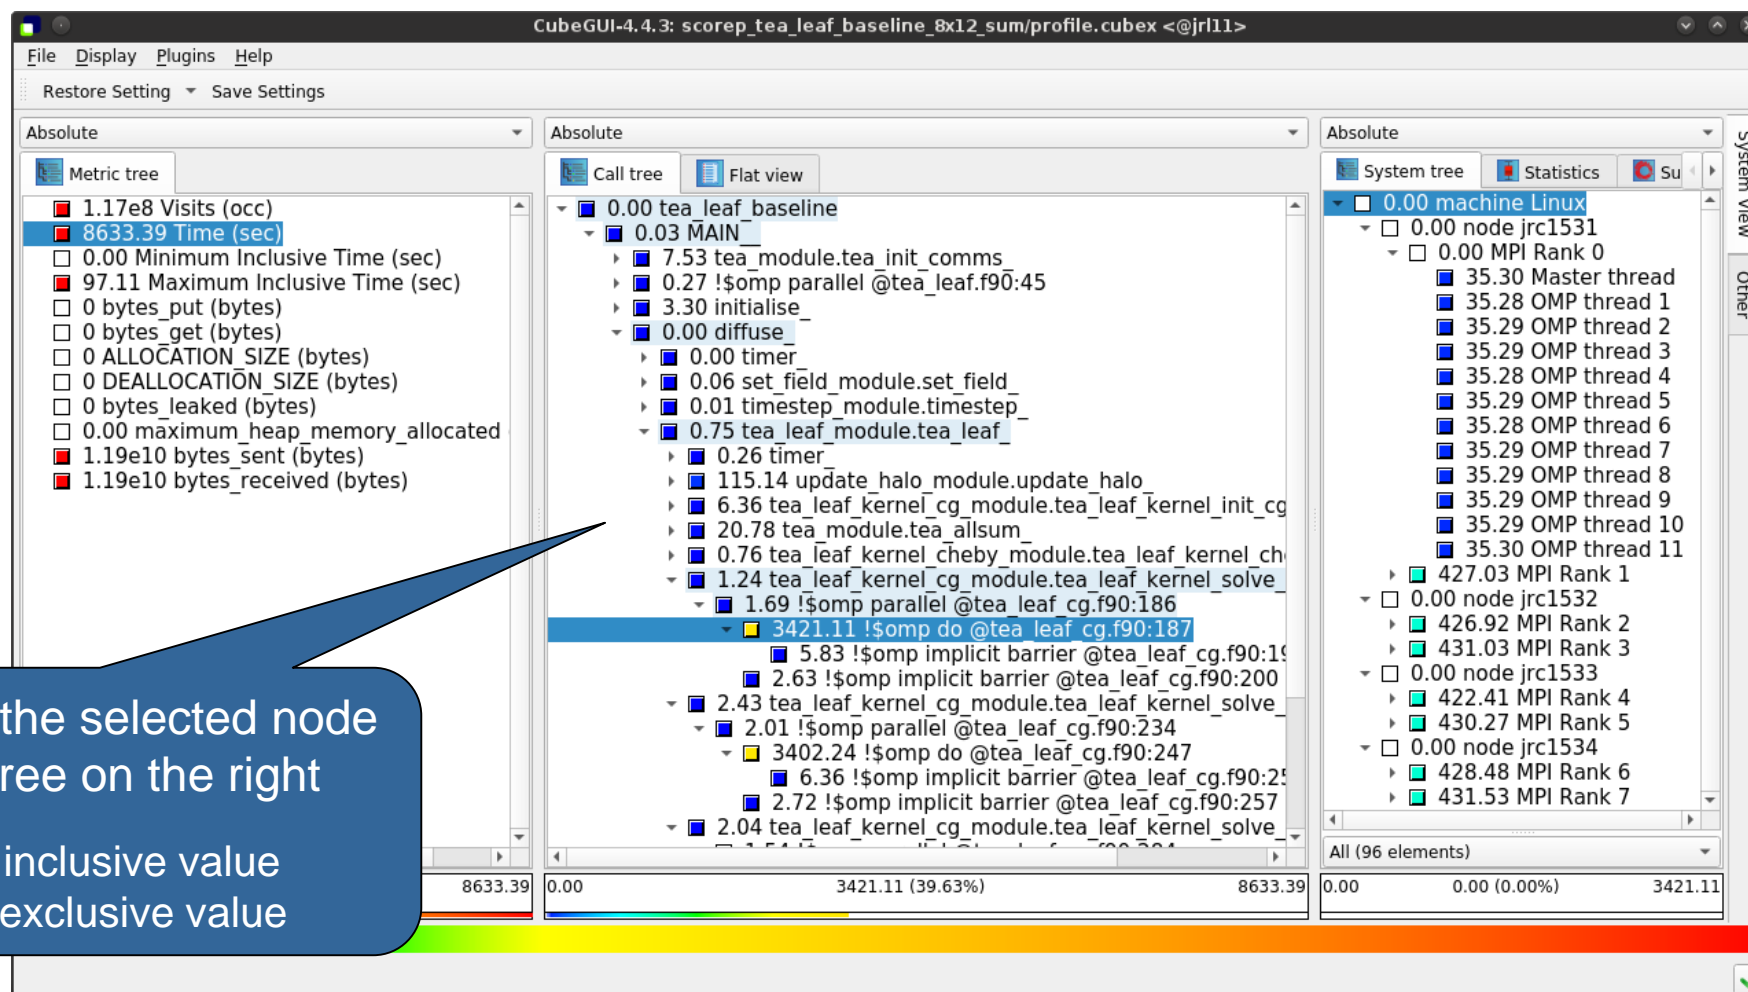


What kind of performance metric?

Where is it in the source code?
In what context?

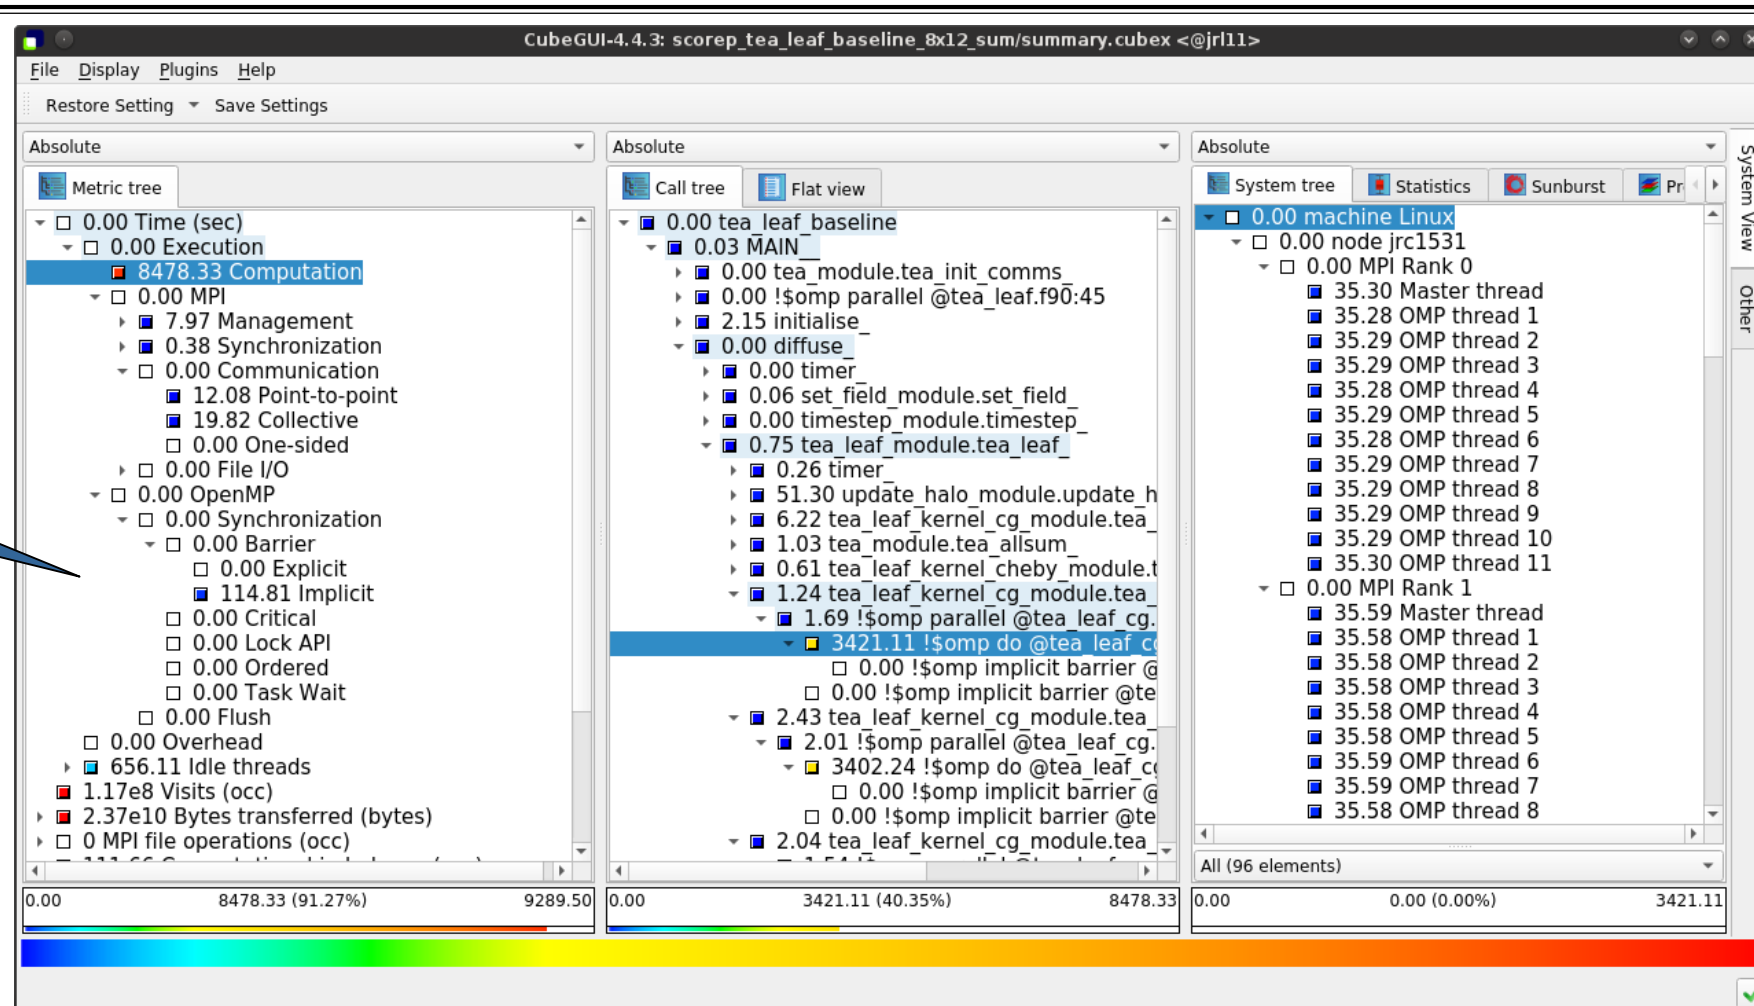
How is it distributed across the processes/threads?

Plain summary analysis report (expanded call tree/system tree)



Post-processed summary analysis report (Scalasca)

Split base metrics from plain report into hierarchy of more specific metrics

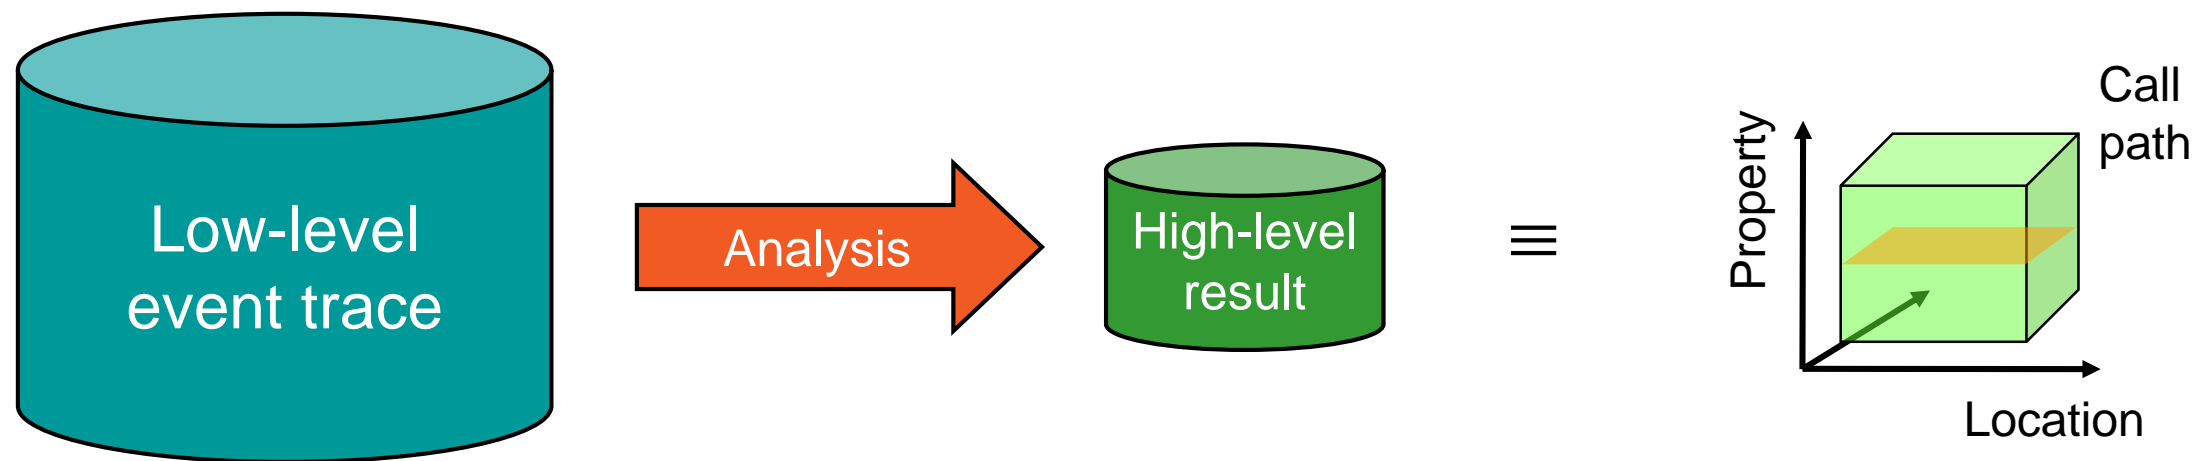


- **Scalable trace-based** performance analysis toolset for the most popular parallel programming paradigms
 - Current focus: MPI, OpenMP, and (to a limited extent) POSIX threads
 - Analysis of traces including only host-side events from applications using CUDA, OpenCL, or OpenACC (also in combination with MPI and/or OpenMP) is possible, but results need to be interpreted with some care
- Specifically targeting large-scale parallel applications
 - Demonstrated scalability up to 1.8 million parallel threads
 - Of course also works at small/medium scale
- Latest release:
 - Scalasca Trace Tools v2.6.1 (Dec 2022)

Automatic trace analysis

▪ Idea

- Automatic search for patterns of inefficient behavior
- Classification of behavior & quantification of significance
- Identification of delays as root causes of inefficiencies



- Guaranteed to cover the entire event trace
- Quicker than manual/visual trace analysis
- Parallel replay analysis exploits available memory & processors to deliver scalability

Scalasca Trace Tools features

- Open source: 3-clause BSD license
- Portability: supports all major HPC platforms
- Scalability: successful analyses with >1M threads
- Uses Score-P instrumenter & measurement libraries
 - Scalasca v2 core package focuses on trace-based analyses
 - Provides convenience commands for measurement, analysis, and postprocessing
 - Supports common data formats
 - Reads event traces in OTF2 format
 - Writes analysis reports in CUBE4 format
- Current limitations:
 - Unable to handle traces ...
 - with MPI thread level exceeding MPI_THREAD_FUNNELED
 - containing memory events, CUDA/HIP/OpenCL device events (kernel, memcpy), SHMEM, or OpenMP nested parallelism
 - PAPI/rusage metrics for trace events are ignored

Putting it all together

