

# Score-P instrumentation and measurement infrastructure

**Demo/Hands-on: Filtering & optimized measurement**

---



# Congratulations!?

---

- If you made it this far, you successfully used Score-P to
  - instrument the application
  - analyze its execution with a summary measurement, and
  - examine it with one of the interactive analysis report explorer GUIs
- ... revealing the call-path profile annotated with
  - the “Time” metric
  - Visit counts
  - MPI message statistics (bytes sent/received)
- ... but how **good** was the measurement?
  - The measured execution produced the desired valid result
  - however, the execution took rather longer than expected!
    - even when ignoring measurement start-up/completion, therefore
    - it was probably dilated by instrumentation/measurement overhead

# Performance analysis steps

---

- 0.0 Reference preparation for validation
- 1.0 Program instrumentation
  - 1.1 Summary measurement collection
  - 1.2 Summary analysis report examination
- 2.0 Summary experiment scoring
  - 2.1 Summary measurement collection with filtering
  - 2.2 Filtered summary analysis report examination
- 3.0 Event trace collection
  - 3.1 Event trace analysis & report examination

# BT-MZ summary analysis result scoring

```
% scorep-score scorep_bt-mz_sum/profile.cubex
```

Estimated aggregate size of event trace:

Estimated requirements for largest trace buffer (max\_buf):

Estimated memory requirements (SCOREP\_TOTAL\_MEMORY):

(warning: The memory requirements cannot be satisfied by Score-P to avoid intermediate flushes when tracing. Set SCOREP\_TOTAL\_MEMORY=4G to get the maximum supported memory or reduce requirements using USR regions filters.)

flt	type	max_buf[B]	visits	time[s]	time[%]	time/visit[us]	region
	ALL	21,395,581,557	6,554,106,209	2340.70	100.0	0.36	ALL
	USR	21,309,225,312	6,537,020,537	1098.88	46.9	0.17	USR
	OMP	83,713,600	16,327,168	1218.52	52.1	74.63	OMP
	COM	2,355,080	724,640	2.91	0.1	4.01	COM
	MPI	287,524	33,856	20.39	0.9	602.30	MPI
	SCOREP	41	8	0.00	0.0	593.30	SCOREP

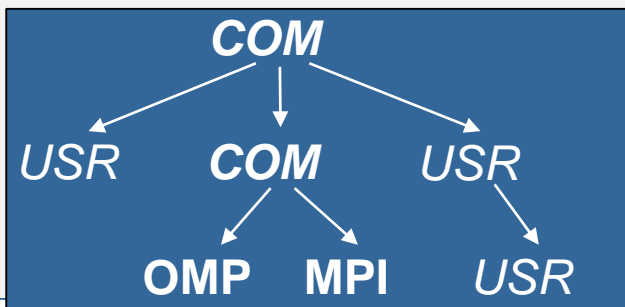
159GB

20GB

20GB

- Report scoring as textual output

~160 GB total memory  
~20 GB per rank!



- Region/callpath classification
  - **MPI** pure MPI functions
  - **OMP** pure OpenMP regions
  - **USR** user-level computation
  - **COM** "combined" USR+OpenMP/MPI
  - **ALL** aggregate of all region types

# BT-MZ summary analysis report breakdown


```
% scorep-score -r scorep_bt-mz_sum/profile.cubex
```

```
[...]
```

```
[...]
```

flt	type	max_buf[B]	visits	time[s]	time[%]	time/visit[us]	region
	ALL	21,395,581,557	6,554,106,209	2340.70	100.0	0.36	ALL
	USR	21,309,225,312	6,537,020,537	1098.88	46.9	0.17	USR
	OMP	83,713,600	16,327,168	1218.52	52.1	74.63	OMP
	COM	2,355,080	724,640	2.91	0.1	4.01	COM
	MPI	287,524	33,856	20.39	0.9	602.30	MPI
	SCOREP	41	8	0.00	0.0	593.30	SCOREP

USR	6,883,222,086	2,110,313,472	422.69	18.1	0.20	binvcrhs_
USR	6,883,222,086	2,110,313,472	284.45	12.2	0.13	matvec_sub_
USR	6,883,222,086	2,110,313,472	360.89	15.4	0.17	matmul_sub_
USR	293,617,584	87,475,200	14.16	0.6	0.16	lhsinit_
USR	293,617,584	87,475,200	11.98	0.5	0.14	binvrhs_
USR	101,320,128	31,129,600	3.68	0.2	0.12	exact_solution_



~20 GB just for these  
6 regions



## BT-MZ summary analysis score

---

- Summary measurement analysis score reveals
  - Total size of event trace would be ~160 GB
  - Maximum trace buffer size would be ~20 GB per rank
    - smaller buffer would require flushes to disk during measurement resulting in substantial perturbation
  - 99.5% of the trace requirements are for USR regions
    - purely computational routines never found on COM call-paths common to communication routines or OpenMP parallel regions
  - These USR regions contribute around 47% of total time
    - however, much of that is very likely to be measurement overhead for frequently-executed small routines
- Advisable to tune measurement configuration
  - Specify an adequate trace buffer size
  - Specify a filter file listing (USR) regions not to be measured

# BT-MZ summary analysis report filtering

```
% cat ../config/scorep.filt
SCOREP_REGION_NAMES_BEGIN
EXCLUDE
  binvcrhs*
  matmul_sub*
  matvec_sub*
  exact_solution*
  binvrhs*
  lhs*init*
  timer_*
SCOREP_REGION_NAMES_END

% scorep-score -f ../config/scorep.filt -c 2 \
  scorep_bt-mz_sum/profile.cubex
```

```
Estimated aggregate size of event trace:
Estimated requirements for largest trace buffer (max_buf):
Estimated memory requirements (SCOREP_TOTAL_MEMORY):
(hint: When tracing set SCOREP_TOTAL_MEMORY=215MB to avoid
       intermediate flushes or reduce requirements using
       USR regions filters.)
```

1621MB  
203MB  
215MB

- Report scoring with prospective filter listing 7 USR regions

1.6 GB of memory in total,  
203 MB per rank!  
(Including 2 metric values)

# BT-MZ summary analysis report filtering

```
% scorep-score -r -f ../config/scorep.filt \
  scorep_bt-mz_sum/profile.cubex
```

flt	type	max_buf[B]	visits	time[s]	time[%]	time/ visit[us]	region
-	ALL	21,395,581,557	6,554,106,209	2340.70	100.0	0.36	ALL
-	USR	21,309,225,312	6,537,020,537	1098.88	46.9	0.17	USR
-	OMP	83,713,600	16,327,168	1218.52	52.1	74.63	OMP
-	COM	2,355,080	724,640	2.91	0.1	4.01	COM
-	MPI	287,524	33,856	20.39	0.9	602.30	MPI
-	SCOREP	41	8	0.00	0.0	593.30	SCOREP
*	ALL	86,356,295	17,085,681	1242.85	53.1	72.74	ALL-FLT
+	FLT	21,309,225,262	6,537,020,528	1097.85	46.9	0.17	FLT
-	OMP	83,713,600	16,327,168	1218.52	52.1	74.63	OMP-FLT
*	COM	2,355,080	724,640	2.91	0.1	4.01	COM-FLT
-	MPI	287,524	33,856	20.39	0.9	602.30	MPI-FLT
*	USR	50	9	1.03	0.0	114496.02	USR-FLT
-	SCOREP	41	8	0.00	0.0	593.30	SCOREP-FLT
+	USR	6,883,222,086	2,110,313,472	422.69	18.1	0.20	binvcrhs_
+	USR	6,883,222,086	2,110,313,472	284.45	12.2	0.13	matvec_sub_
+	USR	6,883,222,086	2,110,313,472	360.89	15.4	0.17	matmul_sub_
+	USR	293,617,584	87,475,200	14.16	0.6	0.16	lhsinit_
+	USR	293,617,584	87,475,200	11.98	0.5	0.14	binvrhs_
+	USR	101,320,128	31,129,600	3.68	0.2	0.12	exact_solution_

- Score report breakdown by region (w/o additional metrics)

Filtered routines marked with '+'



# BT-MZ filtered summary measurement

```
% cd bin.scorep
% cp ../jobscript/dine/scorep.sbatch .
% vim scorep.sbatch

# Score-P measurement configuration
export SCOREP_EXPERIMENT_DIRECTORY=scorep_bt-mz_sum_filter
export SCOREP_FILTERING_FILE=../config/scorep.filt
#export SCOREP_TOTAL_MEMORY=100M
#export SCOREP_METRIC_PAPI=PAPI_TOT_INS,PAPI_TOT_CYC,...
#export SCOREP_ENABLE_TRACING=true

set -x
export OMP_NUM_THREADS=6
time -p mpiexec -np 8 ./bt-mz_C.8

% sbatch scorep.sbatch
```

- Set new experiment directory and re-run measurement with new filter configuration
- Submit job

# Score-P filtering

```
% cat ../config/scorep.filt
SCOREP_REGION_NAMES_BEGIN
EXCLUDE
  binvcrhs*
  matmul_sub*
  matvec_sub*
  exact_solution*
  binvrhs*
  lhs*init*
  timer_*
SCOREP_REGION_NAMES_END

% export SCOREP_FILTERING_FILE=\
../config/scorep.filt
```

Region name  
filter block  
using wildcards

Apply filter

- Filtering by source file name
  - All regions in files that are excluded by the filter are ignored
- Filtering by region name
  - All regions that are excluded by the filter are ignored
  - Overruled by source file filter for excluded files
- Apply filter by
  - exporting `SCOREP_FILTERING_FILE` environment variable
- Apply filter at
  - Run-time
  - Compile-time (GCC-plugin only, Intel in 7.0 release)
    - Add cmd-line option `--instrument-filter`
    - No overhead for filtered regions but recompilation

# Source file name filter block

- Keywords
  - Case-sensitive
  - SCOREP\_FILE\_NAMES\_BEGIN, SCOREP\_FILE\_NAMES\_END
    - Define the source file name filter block
    - Block contains EXCLUDE, INCLUDE rules
  - EXCLUDE, INCLUDE rules
    - Followed by one or multiple white-space separated source file names
    - Names can contain bash-like wildcards \*, ?, []
    - Unlike bash, \* may match a string that contains slashes
  - EXCLUDE, INCLUDE rules are applied in sequential order
  - Regions in source files that are excluded after all rules are evaluated, get filtered

```
# This is a comment
SCOREP_FILE_NAMES_BEGIN
    # by default, everything is included
    EXCLUDE */foo/bar*
    INCLUDE */filter_test.c
SCOREP_FILE_NAMES_END
```

# Region name filter block

- Keywords
  - Case-sensitive
  - SCOREP\_REGION\_NAMES\_BEGIN,  
SCOREP\_REGION\_NAMES\_END
    - Define the region name filter block
    - Block contains EXCLUDE, INCLUDE rules
  - EXCLUDE, INCLUDE rules
    - Followed by one or multiple white-space separated region names
    - Names can contain bash-like wildcards \*, ?, []
  - EXCLUDE, INCLUDE rules are applied in sequential order
  - Regions that are excluded after all rules are evaluated, get filtered

```
# This is a comment
SCOREP_REGION_NAMES_BEGIN
    # by default, everything is included
    EXCLUDE *
    INCLUDE bar foo
           baz
           main
SCOREP_REGION_NAMES_END
```

# Region name filter block, mangling

- Name mangling
  - Filtering based on names seen by the measurement system
    - Dependent on compiler
    - Actual name may be mangled
- scorep-score names as starting point  
(e.g. `matvec_sub_`)
  - Use `*` for Fortran trailing underscore(s) for portability
  - Use `?` and `*` as needed for full signatures or overloading
  - Use `\` to escape special characters

```
void bar(int* a) {  
    *a++;  
}  
int main() {  
    int i = 42;  
    bar(&i);  
    return 0;  
}
```

```
# filter bar:  
# for gcc-plugin, scorep-score  
# displays 'void bar(int*)',  
# other compilers may differ  
  
SCOREP_REGION_NAMES_BEGIN  
    EXCLUDE void?bar(int?)  
SCOREP_REGION_NAMES_END
```