Linaro Forge

HPC Development Solutions from Linaro

Best in class commercially supported tools for Linux and high-performance computing (HPC)

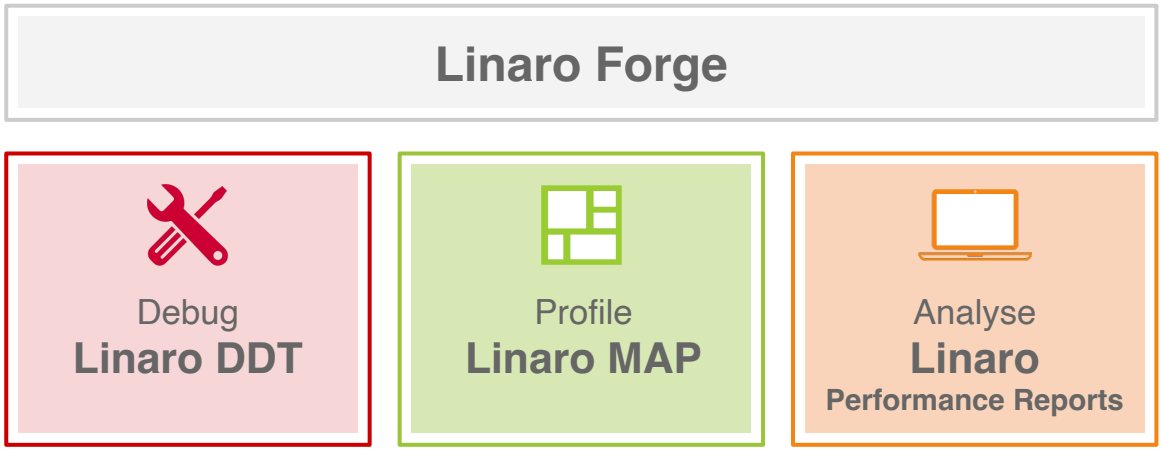

Performance Engineering for any architecture, at any scale

Linaro Forge

An interoperable toolkit for debugging and profiling

The de-facto standard for HPC development

- Most widely-used debugging and profiling suite in HPC
- Fully supported by Linaro on Intel, AMD, Arm, Nvidia, AMD GPUs, etc.

State-of-the art debugging and profiling capabilities

- Powerful and in-depth error detection mechanisms (including memory debugging)
- Sampling-based profiler to identify and understand bottlenecks
- Available at any scale (from serial to exascale applications)

Easy to use by everyone

- Unique capabilities to simplify remote interactive sessions
- Innovative approach to present quintessential information to users

Supported Platforms

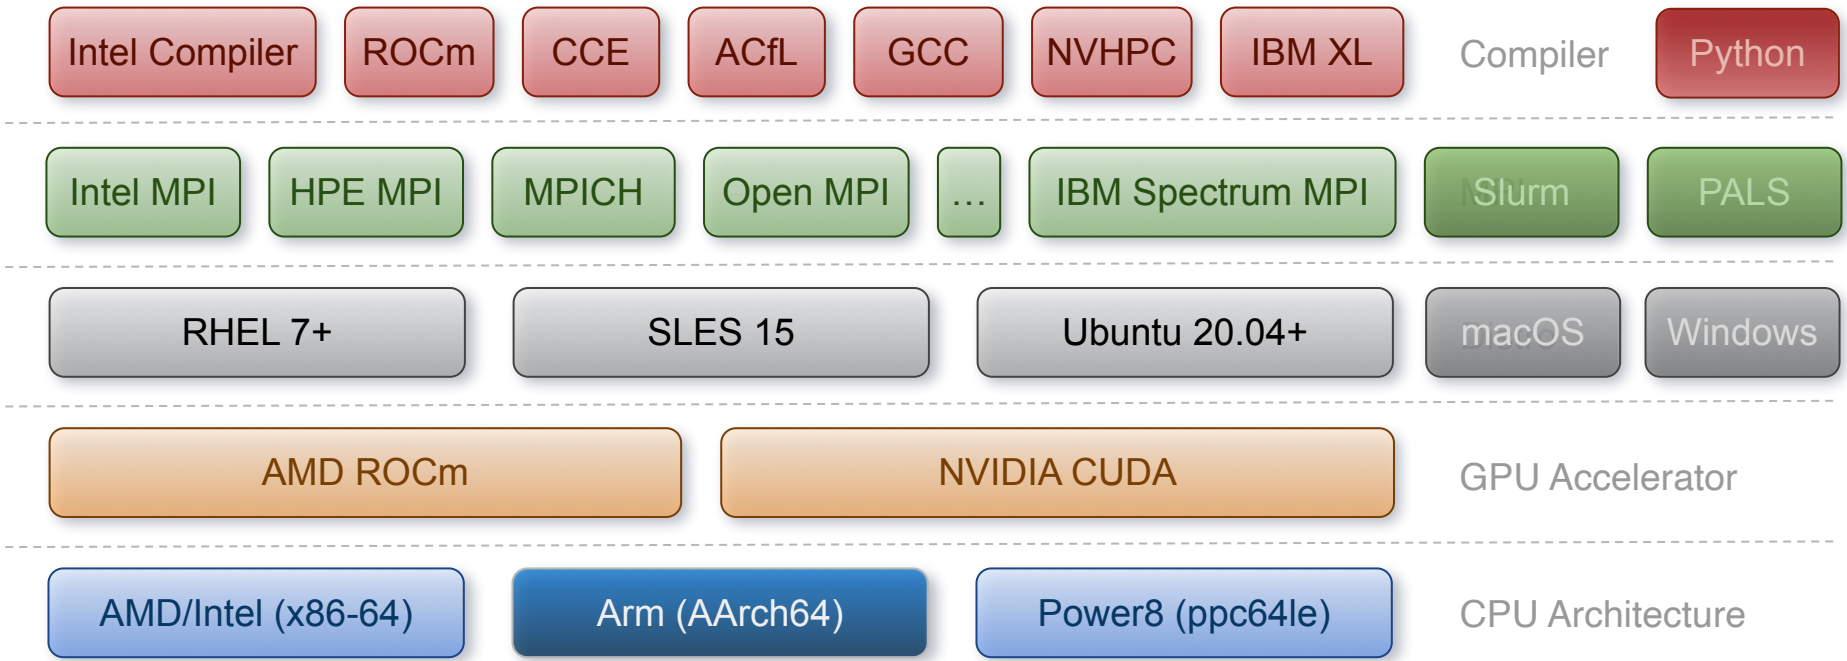

Linaro Performance Reports

A high-level view of application performance with “plain English” insights

Command:	mpiexec.hydra -host node-1,node-2 -map-by socket -n 16 -ppn 8 ./Bin/low_freq/../../Src//hydro -i ./Bin/low_freq/../../../../Input/input_250x125_corner.nml
Resources:	2 nodes (8 physical, 8 logical cores per node)
Memory:	15 GiB per node
Tasks:	16 processes, OMP_NUM_THREADS was 1
Machine:	node-1
Start time:	Thu Jul 9 2015 10:32:13
Total time:	165 seconds (about 3 minutes)
Full path:	Bin/../Src

Summary: hydro is MPI-bound in this configuration

Compute	20.6%	Time spent running application code. High values are usually good. This is **very low**; focus on improving MPI or I/O performance first
MPI	63.2%	Time spent in MPI calls. High values are usually bad. This is **high**; check the MPI breakdown for advice on reducing it
I/O	16.2%	Time spent in filesystem I/O. High values are usually bad. This is **average**; check the I/O breakdown section for optimization advice

I/O

A breakdown of the 16.2% I/O time:

Time in reads	0.0%
Time in writes	100.0%
Effective process read rate	0.00 bytes/s
Effective process write rate	1.38 MB/s

Most of the time is spent in write operations with a very low effective transfer rate. This may be caused by contention for the filesystem or inefficient access patterns. Use an I/O profiler to investigate which write calls are affected.

Linaro Performance Reports Metrics

Lowers expertise requirements by explaining everything in detail right in the report

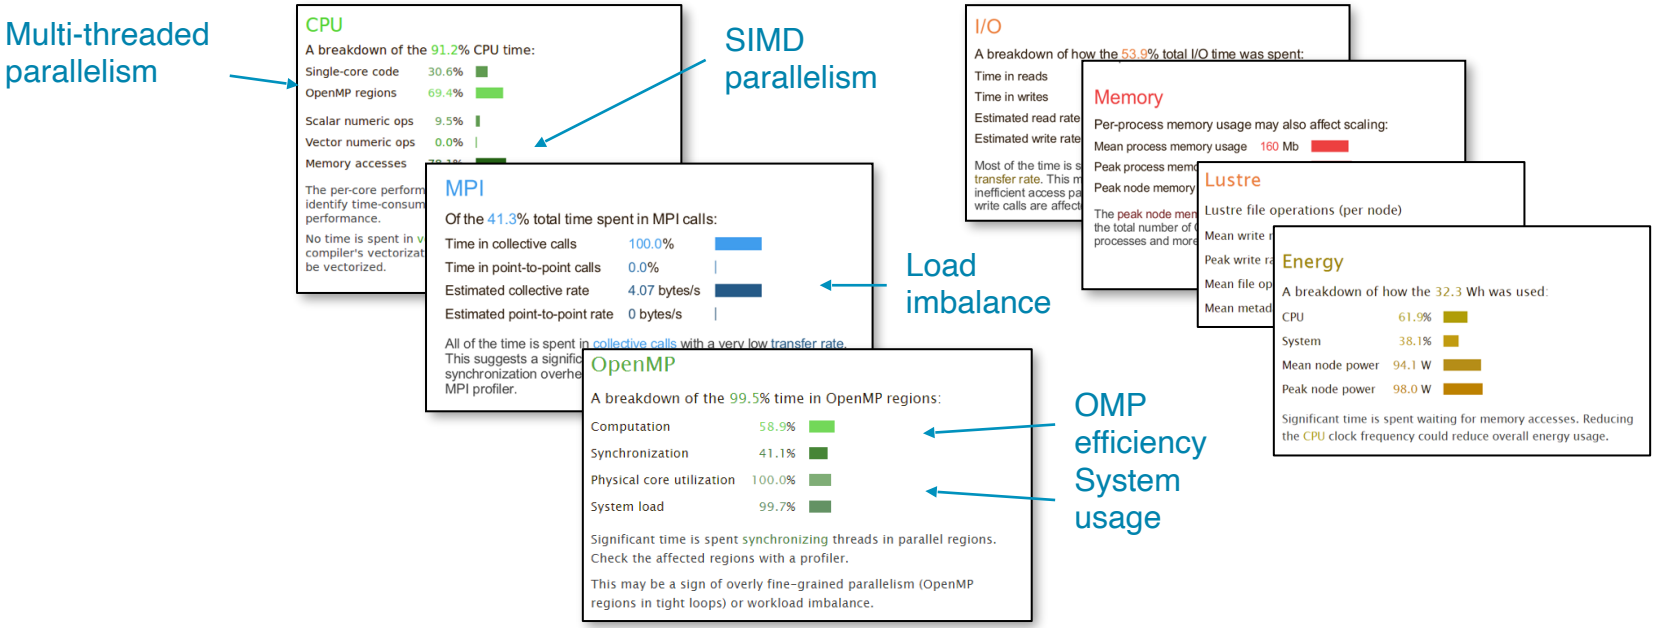

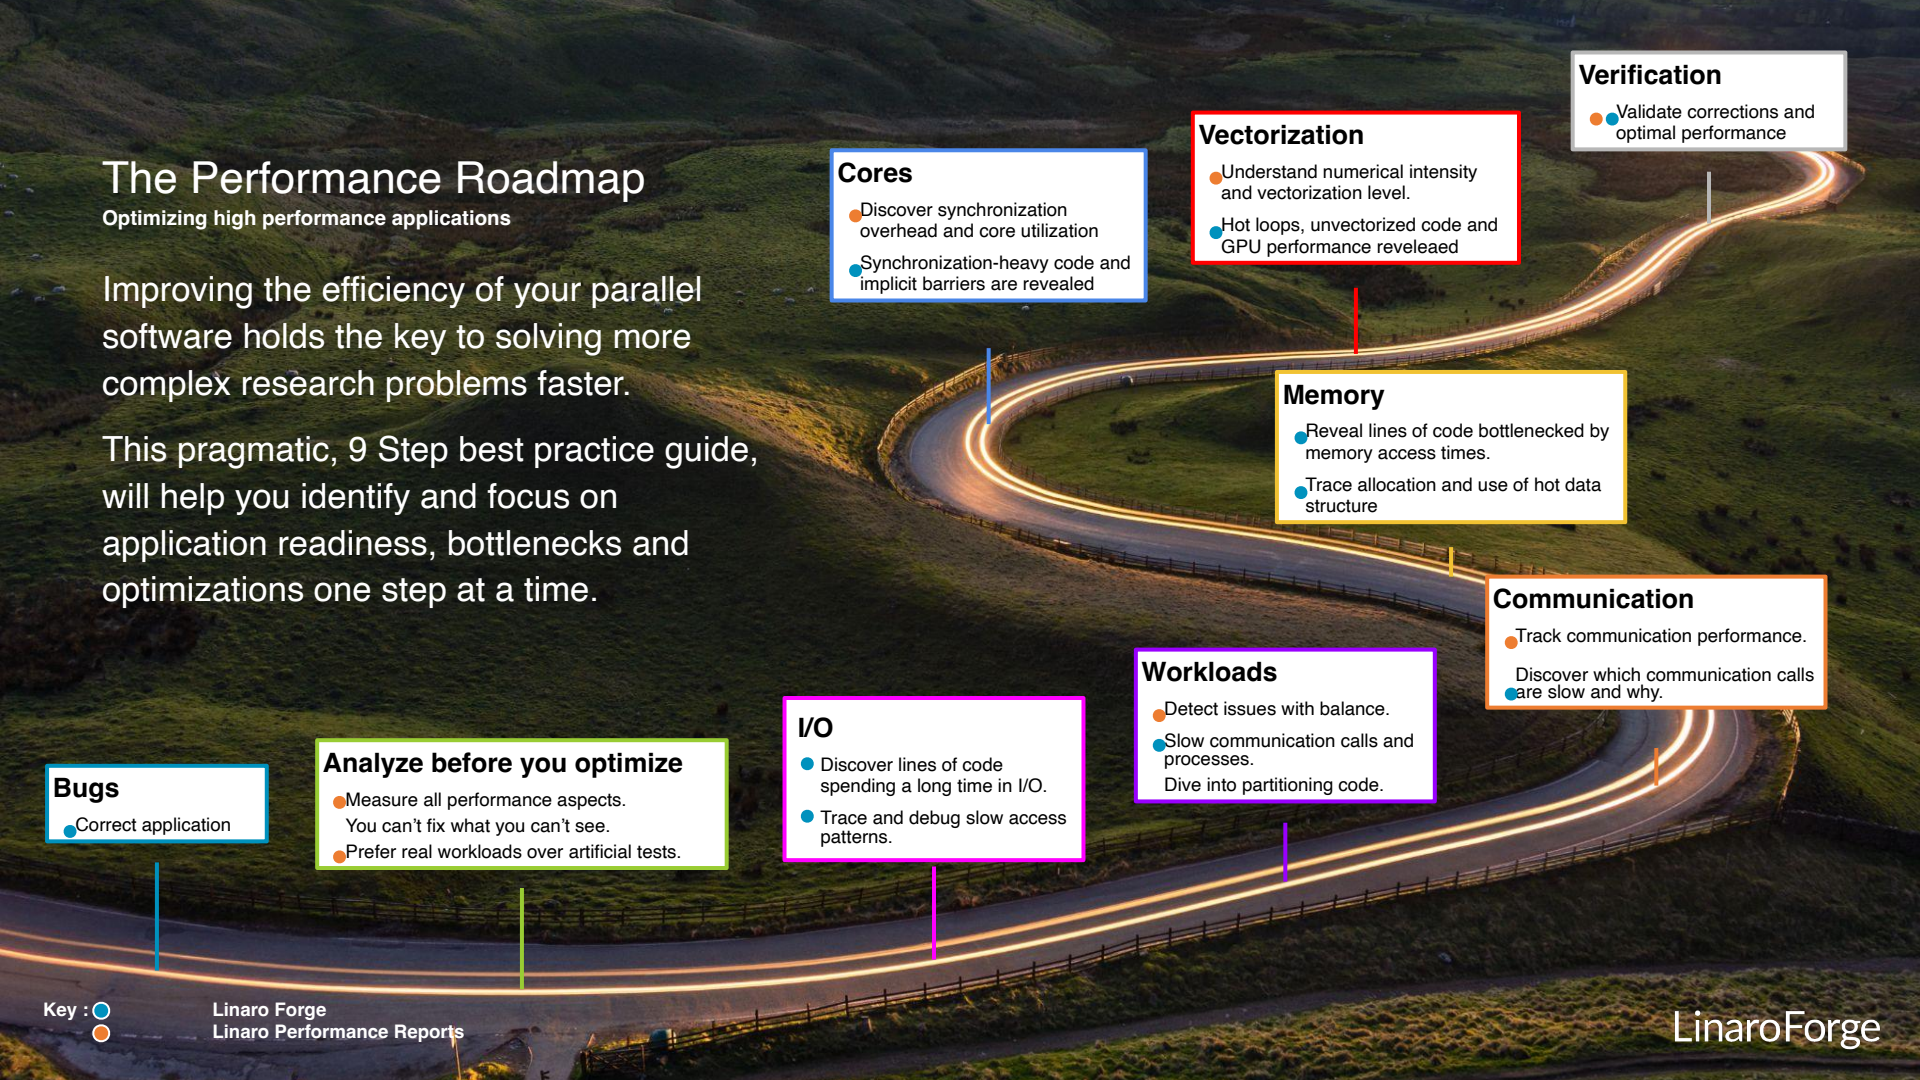
The Performance Roadmap
Optimizing high performance applications
Improving the efficiency of your parallel software holds the key to solving more complex research problems faster.
This pragmatic, 9 Step best practice guide, will help you identify and focus on application readiness, bottlenecks and optimizations one step at a time.
Verification
Validate corrections and optimal performance
Vectorization
Understand numerical intensity and vectorization level.
Hot loops, unvectorized code and GPU performance reveleaed
Cores
Discover synchronization overhead and core utilization
Synchronization-heavy code and implicit barriers are revealed
Memory
Reveal lines of code bottlenecked by memory access times.
Trace allocation and use of hot data structure
Communication
Track communication performance.
Discover which communication calls are slow and why.
Workloads
Detect issues with balance.
Slow communication calls and processes.
Dive into partitioning code.
I/O
Discover lines of code spending a long time in I/O.
Trace and debug slow access patterns.
Analyze before you optimize
Measure all performance aspects.
You can't fix what you can't see.
Prefer real workloads over artificial tests.
Bugs
Correct application
Key :
Linaro Forge
Linaro Performance Reports
LinaroForge

Performance Improvement

i, j, k

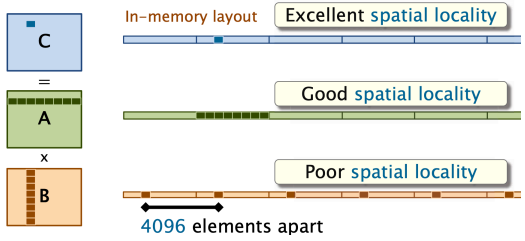

i, k, j

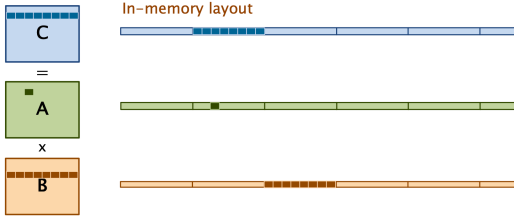

Think,

code,

i, j, k

```
for (int i = 0; i < n; ++i) {
  for (int j = 0; j < n; ++j) {
    for (int k = 0; k < n; ++k) {
      C[i][j] += A[i][k] * B[k][j];
    }
  }
}
```

i, k, j

```
for (int i = 0; i < n; ++i) {
    for (int k = 0; k < n; ++k) {
  for (int j = 0; j < n; ++j) {
      C[i][j] += A[i][k] * B[k][j];
    }
  }
}
```

run, run, run…

…to test and measure many different implementations

Loop order (outer to inner)	Running time (s)
i, j, k	1155.77
i, k, j	177.68
j, i, k	1080.61
j, k, i	3056.63
k, i, j	179.21
k, j, i	3032.82

MAP Capabilities

MAP is a sampling based scalable profiler

- Parallel support for MPI, OpenMP, CUDA
- Designed for C/C++/Fortran

Designed for ‘hot-spot’ analysis

- Stack traces
- Augmented with performance metrics

Adaptive sampling rate

- Throws data away - 1,000 samples per process
- Low overhead, scalable and small file size

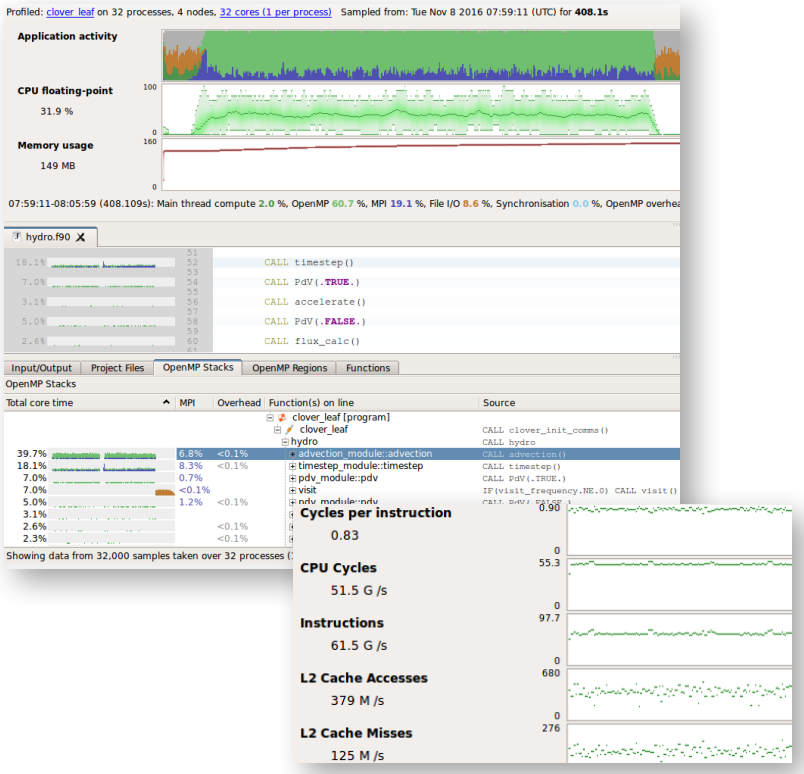

Linaro MAP Source Code Profiler Highlights

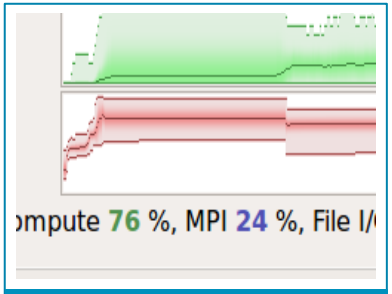

Find the peak memory use

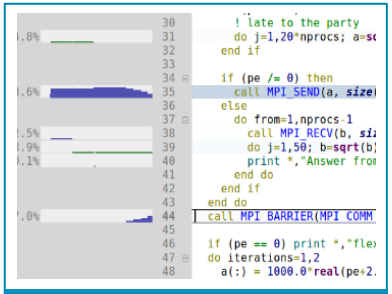

Fix an MPI imbalance

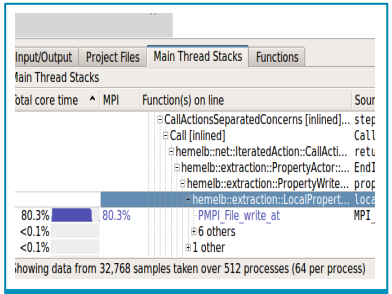

Remove I?O bottleneck

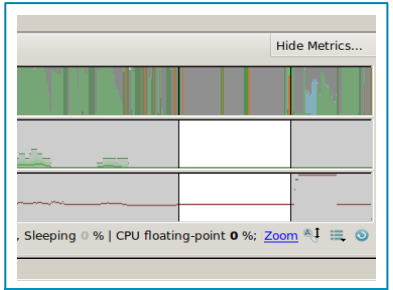

Make sure OpenMP regions make sense

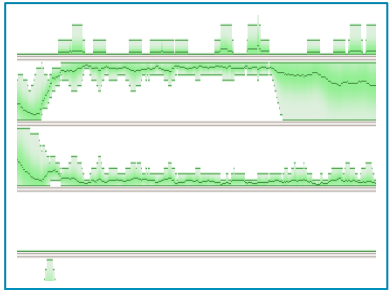

Improve memory access

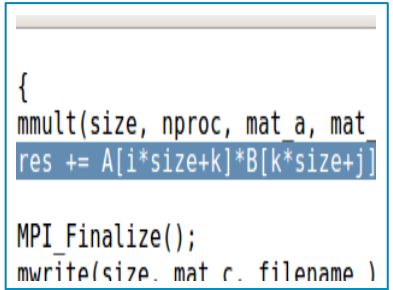

Restructure for vectorization

Toggle percentage-time and core-time in MAP

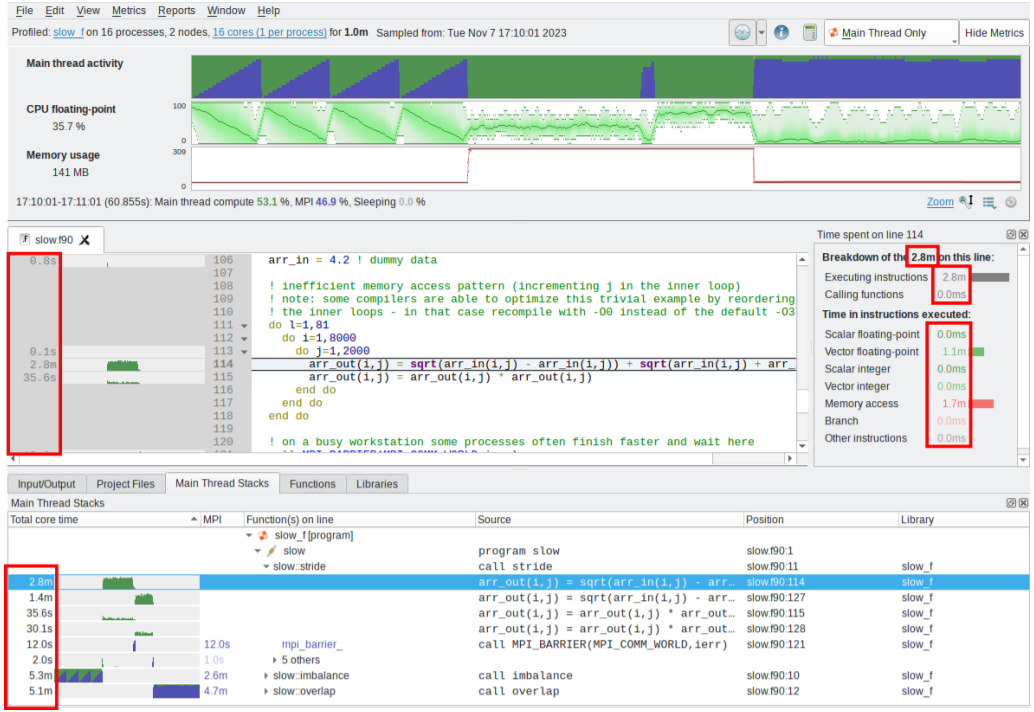

Use for direct comparisons between runs at the same scale (process/core counts).

- Easily determine if a change has made a portion of code faster, slower, or largely unchanged.
- Performance report automatically includes both percentage-time and core time
- Core-time is an estimation, but should be very close to the application run time

Linaro

Hands on Setup

Remote System

Host dine
    Hostname login8.cosma.dur.ac.uk
    user <username>

Examples are in */cosma/home/do009/linaro/performance*

module load allinea/ddt/23.1.0

Local Machine

Install Forge *https://www.linaroforge.com/downloadForge*

Forge userguide

Matrix Multiplication example

Build and run matrix multiplication example

https://docs.linaroforge.com/23.1.1/html/forge/worked_examples_appendix/mmult/analyze.html

Using the Intel modules available on dine

```
# Build C Examples
 make -f mmult.makefile DEBUG=1

# Run with MAP or Performance reports
  map --profile --mpi=generic -n 16 <performance folder>/mmult_c 3072
  perf-report --mpi=generic -n 16 <performance folder>/mmult_c 3072

# Offline profile
 /cosma/home/do009/linaro/submit-job.sh
```

Linaro Forge